

# A PAIRWISE LEARNING-TO-RANK MODEL FOR FEATURE ENGINEERING, LOGISTIC OBJECTIVES, AND CALIBRATION

IZ TECÚM

**ABSTRACT.** We develop an interpretable learning-to-rank pipeline for admissions screening in a small academic society setting (Pi Mu Epsilon), trained from a limited number of human pairwise comparisons. Each applicant  $a$  is mapped deterministically to a feature vector  $x(a) \in \mathbb{R}^d$  derived from structured fields (GPA, calculus completion, advanced coursework), extracted resume text, and short-answer signals. A linear score  $s_w(a) = \langle w, x(a) \rangle$  is learned by minimizing a pairwise logistic loss (Bradley–Terry / RankNet style) with  $\ell_2$  regularization:

$$L_\lambda(w) = \frac{1}{|\mathcal{D}|} \sum_{(i,j,y) \in \mathcal{D}} \log\left(1 + \exp(-z_{ij}\langle w, x(a_i) - x(a_j) \rangle)\right) + \frac{\lambda}{2} \|w\|_2^2, \quad z_{ij} = 2y - 1 \in \{-1, +1\}.$$

Ultimately, we describe the calibration of raw scores to a 0–10 percentile scale and the precise data artifacts that connect deployed scoring to local training (`applicants.jsonl`, `pairs.csv`, and `rank_model.json`).

## CONTENTS

1. Introduction	1
2. Preliminaries	2
3. From applications to vectors	2
3.1. Structured signals from typed fields	3
3.2. Resume extraction and length signal	3
3.3. Keyword detectors as separate coordinates	4
3.4. Short-answer signals and decomposed essay coordinates	4
4. Linear scoring model	5
5. Pairwise training objective	6
5.1. Bradley–Terry / logistic model	6
5.2. Negative log-likelihood loss	6
5.3. Optimization	7
6. Calibration and reporting	7
6.1. Percentile-to-0–10 mapping	7
6.2. Keyword-only summary and decomposition reporting	8
7. Conclusion	8
Acknowledgements	9
References	9

## 1. INTRODUCTION

The operational goal of first-round admissions screening is not to make a final decision, but to produce a consistent, auditable ordering that helps reviewers allocate scarce attention. In our

setting, the data regime is intrinsically small: we may have only  $n \approx 10\text{--}50$  applicants per cycle, and we prefer not to encode “ground-truth” scores that could reflect arbitrary scaling differences across reviewers. This motivates a *pairwise* labeling protocol: reviewers answer questions of the form “between applicant  $i$  and applicant  $j$ , who should rank higher?”

Mathematically, pairwise labels constrain only relative ordering. Let  $\mathcal{A}$  denote the space of applications and let  $x : \mathcal{A} \rightarrow \mathbb{R}^d$  be a deterministic feature map. We seek a score  $s : \mathcal{A} \rightarrow \mathbb{R}$  such that  $s(a_i) > s(a_j)$  whenever the reviewer preference is  $a_i \succ a_j$ . We restrict to linear scores

$$s_w(a) = \langle w, x(a) \rangle, \quad w \in \mathbb{R}^d,$$

to preserve interpretability: each coordinate of  $x(a)$  corresponds to a named signal, and each weight  $w_k$  can be inspected to understand how that signal influences rank.

The training data are a multiset of labeled comparisons

$$\mathcal{D} = \{(i, j, y_{ij})\}, \quad y_{ij} \in \{0, 1\},$$

where  $y_{ij} = 1$  indicates  $a_i \succ a_j$ . A canonical probabilistic model (Bradley–Terry) assumes

$$\mathbb{P}(a_i \succ a_j \mid w) = \sigma(\Delta_{ij}(w)), \quad \Delta_{ij}(w) = s_w(a_i) - s_w(a_j), \quad \sigma(t) = \frac{1}{1 + e^{-t}}.$$

Maximizing likelihood is equivalent to minimizing a sum of logistic losses over score differences, optionally with  $\ell_2$  regularization to stabilize fitting under limited pairs. This setup is widely used in learning-to-rank because it directly matches the type of supervision one can reliably collect (comparisons rather than absolute ratings).

A second design goal is *decomposability*. In particular, we treat resume keyword detectors as separate coordinates (research, teaching, leadership, awards) rather than compressing them into a single opaque number.

## 2. PRELIMINARIES

Let  $\mathcal{A}$  denote the space of applications (structured form fields, course selections, resume PDF text, and short-answer text). Our goal is to output a scalar *ranking score* which is used in an initial screen, while deferring final decisions to human review. Pairwise logistic ranking models are classical in psychometrics and ranking (Bradley–Terry; Plackett–Luce) [1, 3, 2] and are widely used in information retrieval / learning-to-rank [5, 4].

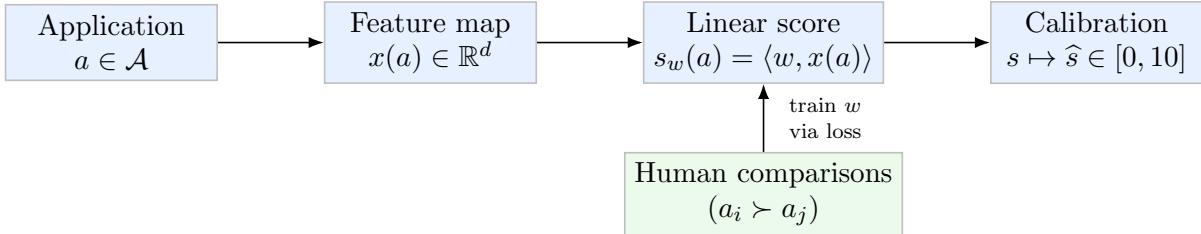


FIGURE 1. Pipeline: applications are mapped to features, scored linearly, calibrated to a 0–10 scale, and trained using pairwise labels.

## 3. FROM APPLICATIONS TO VECTORS

We fix a space of applications  $\mathcal{A}$  and define a *deterministic* feature map

$$x : \mathcal{A} \rightarrow \mathbb{R}^d, \quad a \mapsto x(a) = (x_1(a), \dots, x_d(a)),$$

together with a schema identifier  $\text{fv} \in \mathbb{Z}_{\geq 1}$  (“feature\_version”). The determinism requirement is crucial: for any fixed application  $a$ , repeated evaluation yields the same  $x(a)$ ; all randomness in

the system enters only through the subsequent learning stage (pairwise training) and not through extraction.

The coordinates are designed to be *interpretable*: each  $x_k(a)$  is intended to represent one semantically coherent signal (GPA strength, calculus completion, keyword evidence of research, etc.). In the current implementation we take  $d = 21$  and write the feature vector in blocks,

$$x(a) = (x_{\text{struct}}(a), x_{\text{resume}}(a), x_{\text{kw}}(a), x_{\text{essay}}(a), x_{\text{pen}}(a)),$$

where each block is a small collection of coordinates described below.

**3.1. Structured signals from typed fields.** Let  $\text{GPA}(a) \in [0, 4.33]$  be the reported cumulative GPA (we allow  $A^+$  inflation up to 4.33), let  $\text{Calc}(a) \in \{\text{yes}, \text{no}\}$  denote completion of Calculus I&II (or equivalent), and let  $C(a) \subset \mathcal{C}$  be the set of marked “advanced/recent” courses (with  $|C(a)| \leq 6$ ), validated against a whitelist  $\mathcal{W} \subset \mathcal{C}$  with an explicit exception for graduate courses MATH GR5xxx/6xxx.

We convert these fields into bounded numeric coordinates by explicit monotone maps. First, the GPA coordinate is

$$x_{\text{gpa}}(a) := 10 \cdot \phi\left(\frac{\text{GPA}(a)}{4.33}\right) \in [0, 10],$$

where  $\phi : [0, 1] \rightarrow [0, 1]$  is nondecreasing and chosen to be mildly concave near the top to avoid over-separating near-perfect GPAs. (For instance one may take  $\phi(u) = u^\alpha$  with  $\alpha \in (0, 1]$ , or a piecewise-linear map with a soft plateau; the only properties used later are boundedness and monotonicity.)

Second, calculus completion is encoded as a binary gate

$$x_{\text{calc}}(a) := 10 \cdot \mathbf{1}\{\text{Calc}(a) = \text{yes}\} \in \{0, 10\}.$$

Third, we record a coursework signal that depends only on  $C(a)$

$$x_{\text{upper}}(a) := 10 \cdot \psi\left(\frac{|C(a)|}{6}\right) \in [0, 10],$$

where  $\psi : [0, 1] \rightarrow [0, 1]$  is nondecreasing and may be chosen concave to model diminishing returns from each additional marked course. The validation layer is logically separate: the set  $C(a)$  is constructed by the rule

$$C(a) = \{c \in \mathcal{C} : c \text{ entered by user and } (c \in \mathcal{W} \text{ or } c \text{ is MATH GR5/6xxx})\},$$

so that the score is computed only from accepted courses and is therefore well-defined independent of user formatting noise.

**3.2. Resume extraction and length signal.** Let the applicant upload a PDF resume, from which we extract text in-browser via PDF.js. Denote by  $T(a) \in \Sigma^*$  the extracted plaintext string, and let  $\ell(a) := |T(a)|$  be its character length (or a token length; any fixed notion works provided it is deterministic). We record a length-based completeness proxy

$$x_{\text{rlen}}(a) := 10 \cdot s\left(\frac{\ell(a)}{L_0}\right) \in [0, 10],$$

where  $L_0 > 0$  is a normalization constant and  $s : [0, \infty) \rightarrow [0, 1]$  is a saturating function such as  $s(u) = \min\{1, u\}$  or  $s(u) = 1 - e^{-u}$ . This coordinate is explicitly *not* meant to reward excess word usage; its intended role is to prevent pathological extraction failures (e.g. an empty  $T(a)$  caused by a corrupted upload) from silently propagating through downstream keyword features.

**3.3. Keyword detectors as separate coordinates.** A central design decision is to keep resume keyword evidence in *separate coordinates* rather than collapsing it to one opaque number. Fix four categories

$$\mathcal{K} = \{\text{research, teach, leader, awards}\},$$

and for each  $c \in \mathcal{K}$  fix a finite dictionary  $D_c \subset \Sigma^*$  of tokens/phrases (or regex families) associated with that category. Define the indicator-count statistic

$$N_c(a) := \sum_{w \in D_c} \mathbf{1}\{w \text{ appears in } T(a)\},$$

and map it to a bounded score by a saturating transform with threshold  $\tau_c \geq 1$ ,

$$x_c(a) := 10 \cdot \min\left\{1, \frac{N_c(a)}{\tau_c}\right\} \in [0, 10].$$

Equivalently one may use a smooth saturation such as  $x_c(a) = 10(1 - e^{-N_c(a)/\tau_c})$ ; since  $x$  is deterministic, smoothness is not required for computation, but it can be convenient when reasoning about sensitivity to small textual edits.

This yields four explicit coordinates

$$x_{\text{research}}(a), \quad x_{\text{teach}}(a), \quad x_{\text{leader}}(a), \quad x_{\text{awards}}(a),$$

which the ranker may weight independently. The separability also lets us report a transparent keyword-only summary

$$\text{KW}(a) := \frac{1}{4} \sum_{c \in \mathcal{K}} x_c(a) \in [0, 10],$$

as an auxiliary diagnostic *without changing* the learned score  $s_w(a)$ . In particular, a reviewer can inspect  $\text{KW}(a)$  as a coarse “resume-evidence” meter even if the learned model downweights some keyword categories.

**Example 3.1** (Keyword saturation). Suppose  $D_{\text{research}}$  contains 15 phrases and we set  $\tau_{\text{research}} = 3$ . If an extracted resume contains {“REU”, “preprint”, “poster”} among those phrases, then  $N_{\text{research}}(a) = 3$  and hence  $x_{\text{research}}(a) = 10$ . If it contains only {“REU”}, then  $N_{\text{research}}(a) = 1$  and  $x_{\text{research}}(a) = 10/3 \approx 3.33$ . Thus the coordinate saturates quickly once there is credible evidence, and additional occurrences do not inflate the score beyond the cap.

**3.4. Short-answer signals and decomposed essay coordinates.** Let  $E_A(a), E_B(a) \in \Sigma^*$  denote the two short answers (Short Answer A and Short Answer B). In line with modern ranking practice, we do not treat the essays as a single scalar “quality” label; instead we define several coordinates that measure distinct aspects of writing, mathematical content, and specificity. Formally, we construct a vector

$$x_{\text{essay}}(a) = (x_{\text{band},A}(a), x_{\text{band},B}(a), x_{\text{arc}}(a), x_{\text{plan}}(a), x_{\text{math},A}(a), x_{\text{math},B}(a)) \in [0, 10]^6,$$

where each coordinate is a bounded deterministic functional of  $(E_A(a), E_B(a))$ .

To make this explicit, let  $\text{wc}(E)$  be a word-count functional and  $\text{sent}(E)$  a sentence-count functional. Let  $I_{[2,4]}$  denote the indicator of belonging to the target band. Then a bandpass coordinate can be defined as

$$x_{\text{band},A}(a) := 10 \cdot I_{[2,4]}(\text{sent}(E_A(a))), \quad x_{\text{band},B}(a) := 10 \cdot I_{[2,4]}(\text{sent}(E_B(a))),$$

or, if one prefers robustness to sentence-tokenization artifacts, by a soft band on word counts:

$$x_{\text{band},A}(a) := 10 \cdot \mathbf{1}\{m \leq \text{wc}(E_A(a)) \leq M\},$$

for chosen bounds  $(m, M)$  compatible with the prompt.

For “personal arc” we fix a small dictionary  $R \subset \Sigma^*$  of reflection cues (e.g. *realized, learned, mistake, revised, iterated, feedback, persist*) and define

$$x_{\text{arc}}(a) := 10 \cdot \min \left\{ 1, \frac{1}{\rho} \sum_{w \in R} \mathbf{1}\{w \text{ appears in } E_A(a)\} \right\},$$

with a threshold  $\rho$ . For “plan specificity” in the contribution answer we fix a dictionary  $P$  of concrete action tokens (e.g. *talk, problem session, reading group, outreach, workshop*) and analogously set

$$x_{\text{plan}}(a) := 10 \cdot \min \left\{ 1, \frac{1}{\pi} \sum_{w \in P} \mathbf{1}\{w \text{ appears in } E_B(a)\} \right\}.$$

Finally, math-term density is computed relative to a math lexicon  $M \subset \Sigma^*$  and can be normalized either by word count or by a fixed scale:

$$\begin{aligned} x_{\text{math},A}(a) &:= 10 \cdot \min \left\{ 1, \frac{1}{\mu} \sum_{w \in M} \mathbf{1}\{w \text{ appears in } E_A(a)\} \right\}, \\ x_{\text{math},B}(a) &:= 10 \cdot \min \left\{ 1, \frac{1}{\mu} \sum_{w \in M} \mathbf{1}\{w \text{ appears in } E_B(a)\} \right\}. \end{aligned}$$

The point is structural: each coordinate is an explicit, bounded functional of the raw text, and downstream learning can assign different weights to reflection, specificity, and mathematical content. Such decompositions are standard in learning-to-rank contexts where the supervision is noisy and the model must remain interpretable [5, 8].

**Example 3.2** (Essay decomposition). Consider two applicants with similar GPAs and coursework. Applicant 1 writes  $E_B$  that mentions a specific action (“I will run a weekly problem session on analytic number theory and distribute a short handout”); then  $x_{\text{plan}}$  is near its cap even if  $x_{\text{math},B}$  is moderate. Applicant 2 writes a vague  $E_B$  (“I want to contribute by helping people learn”); then  $x_{\text{plan}}$  is small regardless of other strengths. A linear ranker can therefore separate “intent” from “plan clarity” by placing weight on  $x_{\text{plan}}$  without conflating it with math-term signals.

#### 4. LINEAR SCORING MODEL

We assign a score to an application by

$$s_w(a) = \langle w, x(a) \rangle = \sum_{k=1}^d w_k x_k(a), \quad w \in \mathbb{R}^d.$$

The sign and magnitude of  $w_k$  indicates how feature  $k$  affects rank. Because  $x(a)$  is bounded coordinate-wise, weights are directly interpretable.

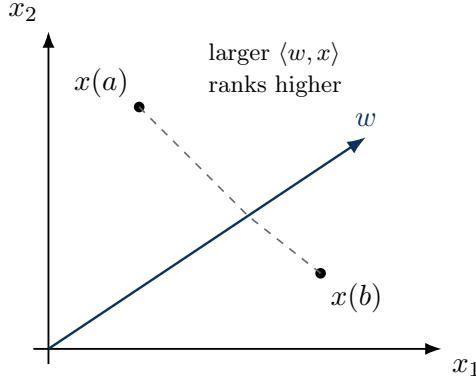


FIGURE 2. Geometric view in  $d = 2$ : applications are points  $x(a)$ ; ranking compares their projections onto  $w$ .

## 5. PAIRWISE TRAINING OBJECTIVE

We obtain training data as labeled pairs. Each label says which of two applications should rank higher.

**Definition 5.1** (Pairwise dataset). A pairwise dataset is a multiset

$$\mathcal{D} = \{(i, j, y_{ij})\},$$

where  $i, j$  index applicants, and  $y_{ij} \in \{0, 1\}$  indicates whether  $i$  should rank above  $j$  ( $y_{ij} = 1$  means  $i \succ j$ ).

**5.1. Bradley–Terry / logistic model.** Define the score difference

$$\Delta_{ij}(w) = s_w(a_i) - s_w(a_j) = \langle w, x(a_i) - x(a_j) \rangle.$$

A standard probabilistic model posits

$$\mathbb{P}(i \succ j \mid w) = \sigma(\Delta_{ij}(w)), \quad \sigma(t) = \frac{1}{1 + e^{-t}}.$$

This is the Bradley–Terry model (in logit form) [1] and is closely related to RankNet [5].

**5.2. Negative log-likelihood loss.** For a single labeled pair  $(i, j, y_{ij})$ , the negative log-likelihood is

$$\ell_{ij}(w) = -y_{ij} \log \sigma(\Delta_{ij}(w)) - (1 - y_{ij}) \log(1 - \sigma(\Delta_{ij}(w))).$$

Equivalently, letting  $z_{ij} = 2y_{ij} - 1 \in \{-1, +1\}$ ,

$$\ell_{ij}(w) = \log(1 + \exp(-z_{ij}\Delta_{ij}(w))),$$

the logistic loss on the signed margin.

We regularize with  $\ell_2$  weight decay:

$$L_\lambda(w) = \frac{1}{|\mathcal{D}|} \sum_{(i,j,y) \in \mathcal{D}} \ell_{ij}(w) + \frac{\lambda}{2} \|w\|_2^2, \quad \lambda \geq 0.$$

**Lemma 5.2** (Gradient). Let  $d_{ij} = x(a_i) - x(a_j)$ . Then

$$\nabla \ell_{ij}(w) = (\sigma(\Delta_{ij}(w)) - y_{ij}) d_{ij},$$

and hence

$$\nabla L_\lambda(w) = \frac{1}{|\mathcal{D}|} \sum_{(i,j,y) \in \mathcal{D}} (\sigma(\langle w, d_{ij} \rangle) - y) d_{ij} + \lambda w.$$

*Proof.* Differentiate  $\ell_{ij}(w)$  using  $\sigma'(t) = \sigma(t)(1 - \sigma(t))$  and the chain rule:

$$\begin{aligned}\frac{\partial}{\partial w}(-y \log \sigma(\Delta)) &= -y \frac{\sigma'(\Delta)}{\sigma(\Delta)} d = -y(1 - \sigma(\Delta)) d, \\ \frac{\partial}{\partial w}(-(1-y) \log(1-\sigma(\Delta))) &= -(1-y) \frac{-\sigma'(\Delta)}{1-\sigma(\Delta)} d = (1-y)\sigma(\Delta) d.\end{aligned}$$

Summing gives  $(\sigma(\Delta) - y)d$ . Adding  $\lambda w$  yields the stated formula.  $\square$

**Proposition 5.3** (Convexity).  $L_\lambda(w)$  is convex in  $w$  for  $\lambda \geq 0$ , and is strictly convex if  $\lambda > 0$ .

*Proof.* Each term  $\ell_{ij}(w) = \log(1 + \exp(-z\langle w, d \rangle))$  is a composition of a convex function  $\log(1 + e^t)$  with an affine map  $t = -z\langle w, d \rangle$ , hence convex. The sum preserves convexity. The quadratic  $\frac{\lambda}{2}\|w\|^2$  is strictly convex when  $\lambda > 0$ , implying strict convexity overall.  $\square$

**5.3. Optimization.** Because  $L_\lambda$  is smooth and convex, standard gradient methods converge reliably [9]. In practice, we run a fixed number of epochs over the dataset and update

$$w^{(t+1)} = w^{(t)} - \eta \nabla L_\lambda(w^{(t)}),$$

with step size  $\eta > 0$  and small  $\lambda$  to reduce overfitting under limited pair counts. This is sufficient at our scale ( $d \approx 21$ ;  $|\mathcal{D}|$  on the order of  $10^1$ – $10^3$ ).

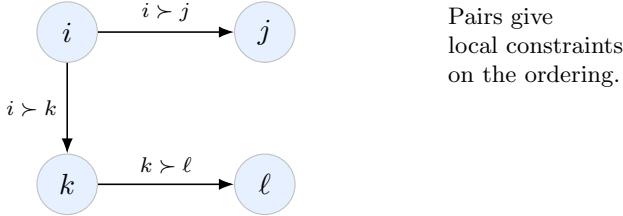


FIGURE 3. Pairwise labels define directed preferences that the model fits in a maximum-likelihood sense.

## 6. CALIBRATION AND REPORTING

The pairwise objective in ?? learns a weight vector  $w \in \mathbb{R}^d$  for the *logit* differences  $s_w(a) - s_w(b) = \langle w, x(a) - x(b) \rangle$ . Consequently, the absolute scale of the raw score

$$s_w(a) = \langle w, x(a) \rangle \in \mathbb{R}$$

is not intrinsically meaningful: multiplying  $w$  by a constant rescales logits and can be partially offset by the regularization and the distribution of labeled pairs. For applicant-facing reporting we therefore apply an explicit *monotone calibration* that preserves order while mapping to a bounded range.

**6.1. Percentile-to-0–10 mapping.** Fix a session (or a reference set) containing a pool of applications  $\mathcal{P} = \{a_1, \dots, a_n\} \subset \mathcal{A}$  and consider the multiset of raw scores  $\{s_w(a_m)\}_{m=1}^n$ . Define the empirical CDF

$$\hat{F}_w(t) := \frac{1}{n} \sum_{m=1}^n \mathbf{1}\{s_w(a_m) \leq t\}, \quad t \in \mathbb{R}.$$

The calibrated score is then the empirical percentile scaled to  $[0, 10]$ ,

$$\hat{s}(a) := 10 \hat{F}_w(s_w(a)) \in [0, 10].$$

This transformation is order-preserving: for any  $a, b \in \mathcal{P}$ ,

$$s_w(a) \leq s_w(b) \implies \hat{s}(a) \leq \hat{s}(b),$$

and it is invariant under strictly increasing reparameterizations of the raw scale. In practice we store both  $s_w(a)$  and  $\hat{s}(a)$ , since the former is the object optimized by the learning-to-rank loss while the latter is the user-facing quantity. When no pool file is available we still log  $s_w(a)$  and may define a degenerate “solo” score  $\hat{s}(a) = 10$  (or  $\hat{s}(a) = 5$ ) purely to keep the interface deterministic; such a score should be treated as non-comparable across sessions.

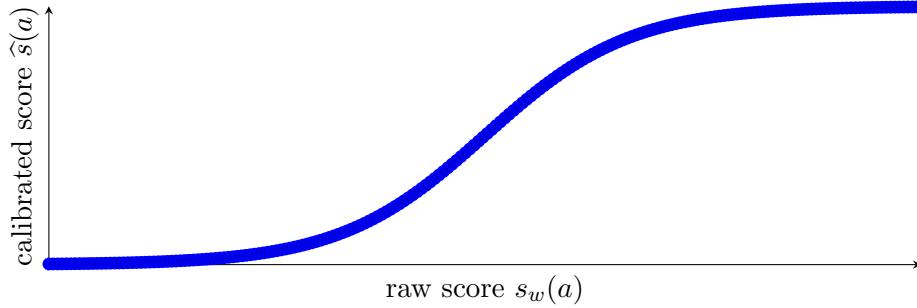


FIGURE 4. A monotone calibration map  $\mathbb{R} \rightarrow [0, 10]$ . In deployment  $\hat{s}(a)$  is computed from the empirical CDF of a pool, i.e. a step function approximation to a CDF; the smooth curve is a schematic proxy.

**6.2. Keyword-only summary and decomposition reporting.** Because the feature map is decomposed into semantically meaningful coordinates, we can expose interpretable slices of  $x(a)$  without altering the trained ranking function. In particular, from Section 3.3 define the keyword-only summary

$$\text{KW}(a) := \frac{1}{4} (x_{\text{research}}(a) + x_{\text{teach}}(a) + x_{\text{leader}}(a) + x_{\text{awards}}(a)) \in [0, 10].$$

This quantity is *not* a replacement for  $s_w(a)$ : it ignores essays, structured eligibility signals, penalties, and any interaction the learned weights impose. Rather,  $\text{KW}(a)$  is a diagnostic scalar that isolates one input subsystem (resume keyword evidence) and can be used in reviewer dashboards, sanity checks, and post-hoc audits. A useful practice is to store the tuple

$$(s_w(a), \hat{s}(a), \text{KW}(a), x(a))$$

in the submission database so that one can answer questions of the form “is this applicant’s high score driven mainly by keywords or by essays/coursework?” without reverse-engineering a monolithic score.

## 7. CONCLUSION

The big thing here is that there is the linear scoring function  $s_w(a) = \langle w, x(a) \rangle$  trained from comparisons via regularized pairwise logistic risk, which yields a convex objective and an interpretable parameterization. The deterministic feature map decomposes an application into bounded, human-legible signals (structured eligibility fields, validated coursework, resume-derived keyword evidence, and short-answer heuristics). Finally, a monotone percentile calibration produces a stable 0–10 reporting score while preserving rank order, and auxiliary summaries such as  $\text{KW}(a)$  provide transparency without entangling the learned objective.

## ACKNOWLEDGEMENTS

I wish recognize my co-founder, Dmytro Prosiannikov, a fellow Columbia physics classmate, for deep discussions that directly shaped the original conception of this ML project and its implementation direction.

## REFERENCES

- [1] R. A. Bradley and M. E. Terry, *Rank analysis of incomplete block designs: I. The method of paired comparisons*, Biometrika **39** (1952), 324–345.
- [2] R. D. Luce, *Individual Choice Behavior: A Theoretical Analysis*, Wiley, 1959.
- [3] R. L. Plackett, *The analysis of permutations*, Journal of the Royal Statistical Society. Series C (Applied Statistics) **24** (1975), 193–202.
- [4] T. Joachims, *Optimizing search engines using clickthrough data*, Proceedings of KDD (2002), 133–142.
- [5] C. J. C. Burges, T. Shaked, E. Renshaw, A. Lazier, M. Deeds, N. Hamilton, and G. Hullender, *Learning to rank using gradient descent*, Proceedings of ICML (2005), 89–96.
- [6] Z. Cao, T. Qin, T.-Y. Liu, M.-F. Tsai, and H. Li, *Learning to rank: From pairwise approach to listwise approach*, Proceedings of ICML (2007), 129–136.
- [7] T.-Y. Liu, *Learning to Rank for Information Retrieval*, Foundations and Trends in Information Retrieval **3** (2009), no. 3, 225–331.
- [8] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning*, 2nd ed., Springer, 2009.
- [9] J. Nocedal and S. J. Wright, *Numerical Optimization*, 2nd ed., Springer, 2006.